Lecture 1: Introduction to the Course Basics of Data

S. Morteza Najibi

Persian Gulf University

February 20, 2017

Introduction

Agenda

▷ Course overview and mechanics
▷ Built-in data types
▷ Built-in functions and operators
▷ First data structures: Vectors and arrays

Why good statisticians learn to program

▷ *Independence*: Otherwise, you rely on someone else having given you exactly the right tool

▷ *Honesty*: Otherwise, you end up distorting your problem to match the tools you have

▷ *Clarity*: Making your method something a machine can do disciplines your thinking and makes it public; that's science

How this class will work

▷ No programming knowledge presumed
▷ Some stats. knowledge presumed
▷ General programming mixed with data-manipulation and statistical inference
▷ Class will be *very* cumulative
▷ Keep up with the readings and assignments!

assignment

Assignments, class notes, grading policies, useful links on R:

▷   `http://cw.sdat.ir`

R Basics

The class in a nutshell: Functional programming

2 sorts of things (**objects**): **data** and **functions**

▷ **Data**: things like 7, "seven", 7.000, the matrix $\begin{bmatrix} 7 & 7 & 7 \\ 7 & 7 & 7 \end{bmatrix}$

▷ **Functions**: things like $\log$, $+$ (two arguments), $<$ (two), mod (two), mean (one)

*A function is a machine which turns input objects (**arguments**) into an output object (**return value**), possibly with **side effects**, according to a definite rule*

Continue

Programming is writing functions to transform inputs into outputs

Good programming ensures the transformation is done easily and correctly

Machines are made out of machines; functions are made out of functions, like $f(a, b) = a^2 + b^2$

> *The route to good programming is to take the big transformation and break it down into smaller ones, and then break those down, until you come to tasks which the built-in functions can do*

Before functions, data

Different kinds of data object

All data is represented in binary format, by **bits** (TRUE/FALSE, YES/NO, 1/0)

▷ **Booleans** Direct binary values: TRUE or FALSE in R
▷ **Integers**: whole numbers (positive, negative or zero), represented by a fixed-length block of bits
▷ **Characters** fixed-length blocks of bits, with special coding; **strings** = sequences of characters
▷ **Floating point numbers**: a fraction (with a finite number of bits) times an exponent, like 1.87×10^6, but in binary form
▷ **Missing or ill-defined values**: NA, NaN, etc.

Operators

▷ **Unary -** for arithmetic negation, **!** for Boolean
▷ **Binary** usual arithmetic operators, plus ones for modulo and integer division;
take two numbers and give a number

some basics

```r
7+5
```

```
## [1] 12
```

```r
7-5
```

```
## [1] 2
```

```r
7*5
```

```
## [1] 35
```

```r
7^5
```

```
## [1] 16807
```

Continue

```r
7/5
```

```
## [1] 1.4
```

```r
7 %% 5
```

```
## [1] 2
```

```r
7 %/% 5
```

```
## [1] 1
```

The R Console

Basic interaction with R is by typing in the **console**, a.k.a. **terminal** or **command-line**

You type in commands, R gives back answers (or errors)

Menus and other graphical interfaces are extras built on top of the console

Operators cont'd.

Comparisons are also binary operators; they take two objects, like numbers, and give a Boolean

```
7 > 5
```

```
## [1] TRUE
```

```
7 < 5
```

```
## [1] FALSE
```

```
7 >= 7
```

```
## [1] TRUE
```

```
7 <= 5
```

```
## [1] FALSE
```

```r
7 == 5
```

```
## [1] FALSE
```

```r
7 != 5
```

```
## [1] TRUE
```

Boolean operators

Basically "and" and "or":

```
(5 > 7) & (6*7 == 42)
```

[1] FALSE

```
(5 > 7) | (6*7 == 42)
```

[1] TRUE

(will see special doubled forms, `&&` and `||`, later)

More types

typeof() function returns the type

is._foo_**()** functions return Booleans for whether the argument is of type _foo_

as._foo_**()** (tries to) "cast" its argument to type _foo_ — to translate it sensibly into a _foo_-type value

```r
typeof(7)
```

```
## [1] "double"
```

```r
is.numeric(7)
```

```
## [1] TRUE
```

```r
is.na(7)
```

```
## [1] FALSE
```

```r
is.na(7/0)
```

```
## [1] FALSE
```

```r
is.na(0/0)
```

```
## [1] TRUE
```

Why is 7/0 not NA, but 0/0 is?

```r
is.character(7)
```

```
## [1] FALSE
```

```r
is.character("7")
```

```
## [1] TRUE
```

```r
is.character("seven")
```

```
## [1] TRUE
```

```r
is.na("seven")
```

```
## [1] FALSE
```

```r
as.character(5/6)
```

```
## [1] "0.833333333333333"
```

```r
as.numeric(as.character(5/6))
```

```
## [1] 0.8333333
```

```r
6*as.numeric(as.character(5/6))
```

```
## [1] 5
```

```r
5/6 == as.numeric(as.character(5/6))
```

```
## [1] FALSE
```

(why is that last FALSE?)

Data can have names

We can give names to data objects; these give us **variables**

A few variables are built in:

```
pi
```

```
## [1] 3.141593
```

Variables can be arguments to functions or operators, just like constants:

```
pi*10
```

```
## [1] 31.41593
```

```
cos(pi)
```

```
## [1] -1
```

Most variables are created with the **assignment operator**, `<-` or `=`

```r
approx.pi <- 22/7
approx.pi
```

```
## [1] 3.142857
```

```r
diameter.in.cubits = 10
approx.pi*diameter.in.cubits
```

```
## [1] 31.42857
```

The assignment operator also changes values:

```
circumference.in.cubits <- approx.pi*diameter.in.cubits
circumference.in.cubits
```

```
## [1] 31.42857
```

```
circumference.in.cubits <- 30
circumference.in.cubits
```

```
## [1] 30
```

Workspace

Using Names

Using names and variables makes code: easier to design, easier to debug, less prone to bugs, easier to improve, and easier for others to read

Avoid "magic constants"; use named variables you will be graded on this!

Named variables are a first step towards **abstraction**

The workspace

What names have you defined values for?

```
ls()
```

```
## [1] "approx.pi"              "circumference.in.cubits"
## [3] "diameter.in.cubits"
```

```
objects()
```

```
## [1] "approx.pi"              "circumference.in.cubits"
## [3] "diameter.in.cubits"
```

Getting of variables:

```
rm("circumference.in.cubits")
ls()
```

```
## [1] "approx.pi"           "diameter.in.cubits"
```

First data structure: vectors

Group related data values into one object, a **data structure**

A **vector** is a sequence of values, all of the same type

```
x <- c(7, 8, 10, 45)
x
```

```
## [1]  7  8 10 45
```

```
is.vector(x)
```

```
## [1] TRUE
```

c() function returns a vector containing all its arguments in order x[1] is the first element, x[4] is the 4th element
x[-4] is a vector containing all but the fourth element

`vector(`length=6`)` returns an empty vector of length 6; helpful for filling things up later

```
weekly.hours <- vector(length=5)
weekly.hours[5] <- 8
```

Vector arithmetic

Operators apply to vectors "pairwise" or "elementwise":

```
y <- c(-7, -8, -10, -45)
x+y
```

```
## [1] 0 0 0 0
```

```
x*y
```

```
## [1]   -49   -64  -100 -2025
```

Recycling

Recycling repeat elements in shorter vector when combined with longer

```
x + c(-7,-8)
```

```
## [1]  0  0  3 37
```

```
x^c(1,0,-1,0.5)
```

```
## [1] 7.000000 1.000000 0.100000 6.708204
```

Single numbers are vectors of length 1 for purposes of recycling:

```
2*x
```

```
## [1] 14 16 20 90
```

Can also do pairwise comparisons:

```
x > 9
```

```
## [1] FALSE FALSE  TRUE  TRUE
```

Note: returns Boolean vector

Boolean operators work elementwise:

```
(x > 9) & (x < 20)
```

```
## [1] FALSE FALSE  TRUE FALSE
```

To compare whole vectors, best to use `identical()` or `all.equal()`:

```
x == -y
```

```
## [1] TRUE TRUE TRUE TRUE
```

```
identical(x,-y)
```

```
## [1] TRUE
```

```
identical(c(0.5-0.3,0.3-0.1),c(0.3-0.1,0.5-0.3))
```

```
## [1] FALSE
```

```
all.equal(c(0.5-0.3,0.3-0.1),c(0.3-0.1,0.5-0.3))
```

```
## [1] TRUE
```

Simple Functions

Functions on vectors

Lots of functions take vectors as arguments: - `mean()`, `median()`, `sd()`, `var()`, `max()`, `min()`, `length()`, `sum()`: return single numbers - `sort()` returns a new vector - `hist()` takes a vector of numbers and produces a histogram, a highly structured object, with the side-effect of making a plot - Similarly `ecdf()` produces a cumulative-density-function object - `summary()` gives a five-number summary of numerical vectors - `any()` and `all()` are useful on Boolean vectors

Addressing vectors

Vector of indices:

```
x[c(2,4)]
```

```
## [1]  8 45
```

Vector of negative indices

```
x[c(-1,-3)]
```

```
## [1]  8 45
```

(why that, and not `8 10`?)

Boolean vector

```
x[x>9]
```

```
## [1] 10 45
```

```
y[x>9]
```

```
## [1] -10 -45
```

which() turns a Boolean vector in vector of TRUE indices:

```
places <- which(x > 9)
places
```

```
## [1] 3 4
```

```
y[places]
```

```
## [1] -10 -45
```

Named components

You can give names to elements or components of vectors

```r
names(x) <- c("v1","v2","v3","fred")
names(x)
```

```
## [1] "v1"    "v2"    "v3"    "fred"
```

```r
x[c("fred","v1")]
```

```
## fred   v1
##   45    7
```

note the labels in what R prints; not actually part of the value

names(x) is just another vector (of characters):

```r
names(y) <- names(x)
sort(names(x))
```

```
## [1] "fred" "v1"    "v2"    "v3"
```

```r
which(names(x)=="fred")
```

```
## [1] 4
```

Take-Aways

▷ We write programs by composing functions to manipulate data
▷ The basic data types let us represent Booleans, numbers, and characters
▷ Data structure let us group related values together
▷ Vectors let us group values of the same type
▷ Use variables rather a profusion of magic constants
▷ Name components of structures to make data more meaningful

Peculiarites of floating-point numbers

The more bits in the fraction part, the more precision

The R floating-point data type is a `double`, a.k.a. `numeric` back when memory was expensive, the now-standard number of bits was twice the default

Finite precision $\Rightarrow$ arithmetic on `doubles` $\neq$ arithmetic on $\mathbb{R}$.

```r
0.45 == 3*0.15
```

```
## [1] FALSE
```

```r
0.45 - 3*0.15
```

```
## [1] 5.551115e-17
```

Often ignorable, but not always - Rounding errors tend to accumulate in long calculations - When results should be ≈ 0, errors can flip signs - Usually better to use `all.equal()` than exact comparison

```
(0.5 - 0.3) == (0.3 - 0.1)
```

```
## [1] FALSE
```

```
all.equal(0.5-0.3, 0.3-0.1)
```

```
## [1] TRUE
```

Peculiarities of Integers

Typing a whole number in the terminal doesn't make an integer; it makes a double, whose fractional part is 0

```
is.integer(7)
```

```
## [1] FALSE
```

This looks like an integer

```
as.integer(7)
```

```
## [1] 7
```

To test for being a whole number, use round():

```
round(7) == 7
```

```
## [1] TRUE
```

Thank You