



Exploratory Data Analysis in R

S. Morteza Najibi, Perisan Gulf University

April 15, 2017

Graphics in R



- ▶ Exploratory Data Analysis (EDA)
- ▶ Powerful environment for visualizing scientific data
- ▶ Fully programmable
- ▶ An example of Exploratory Data Analysis will be performed with the Iris data set.

Important high-level plotting functions



- ▶ `plot`: generic x-y plotting
 - ▶ `barplot`: bar plots
 - ▶ `boxplot`: box-and-whisker plot
 - ▶ `hist`: histograms
 - ▶ `pie`: pie charts
 - ▶ `dotchart`: cleveland dot plots
 - ▶ `image`, `heatmap`, `contour`, `persp`: functions to generate image-like plots
 - ▶ `qqnorm`, `qqline`, `qqplot`: distribution comparison plots
 - ▶ `pairs`, `coplot`, `matplot`: display of multivariant data
- Help on these functions

Iris Data



The iris data set is available at the UCI Machine Learning Repository.

The data set was created by statistician Douglas Fisher. The data consists of 3 classes of flower types:

- ▶ *setosa*
- ▶ *virginica*
- ▶ *versicolor*

The Iris Attributes



The data has 4 attributes:

- ▷ sepal width
- ▷ sepal length
- ▷ petal width
- ▷ petal length

Load Data



The data needs to be loaded into R.

```
head(iris)
```

```
##      Sepal.Length Sepal.Width Petal.Length Petal.Width Species
## 1           5.1           3.5           1.4           0.2   setosa
## 2           4.9           3.0           1.4           0.2   setosa
## 3           4.7           3.2           1.3           0.2   setosa
## 4           4.6           3.1           1.5           0.2   setosa
## 5           5.0           3.6           1.4           0.2   setosa
## 6           5.4           3.9           1.7           0.4   setosa
```

```
names(iris)
```

```
## [1] "Sepal.Length" "Sepal.Width"  "Petal.Length" "Petal.Width"
## [5] "Species"
```

Summary Statistics



The summary function can be used to assess the statistical properties of each attribute.

```
summary(iris[, -5])
```

##	Sepal.Length	Sepal.Width	Petal.Length	Petal.Width
##	Min. :4.300	Min. :2.000	Min. :1.000	Min. :0.100
##	1st Qu.:5.100	1st Qu.:2.800	1st Qu.:1.600	1st Qu.:0.300
##	Median :5.800	Median :3.000	Median :4.350	Median :1.300
##	Mean :5.843	Mean :3.057	Mean :3.758	Mean :1.199
##	3rd Qu.:6.400	3rd Qu.:3.300	3rd Qu.:5.100	3rd Qu.:1.800
##	Max. :7.900	Max. :4.400	Max. :6.900	Max. :2.500

other statistics may be calculated as described in the introduction to R examples.

A decorative background pattern consisting of a grid of triangles and squares. The pattern is composed of light gray lines forming a series of nested and overlapping geometric shapes, creating a complex, crystalline structure. The pattern is centered and extends across the entire slide.


Visualizations

Plot



▷ plot x versus y

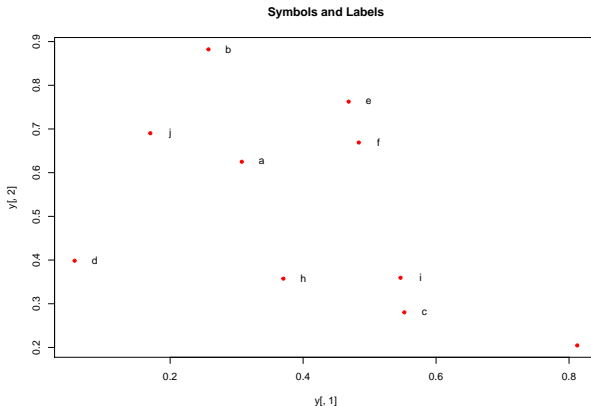
```
plot(iris[,1], iris[,2],xlab=names(iris)[1],ylab=names(iris)[2],col=2)
```



Some more techniques in plot




```
set.seed(100)
y <- matrix(runif(30), ncol=3, dimnames=list(letters[1:10], LETTERS[1:3]))
plot(y[,1], y[,2], pch=20, col="red", main="Symbols and Labels")
text(y[,1]+0.03, y[,2], rownames(y))
```



Add a regression line


```
plot(y[,1], y[,2])  
myline <- lm(y[,2]~y[,1]);  
abline(myline, lwd=2)
```



Add a mathematical expression to a plot



```
plot(y[,1], y[,2])  
text(y[1,1], y[1,2]+.2,expression(sum(frac(1,sqrt(x^2*pi)))), cex=1.3)
```

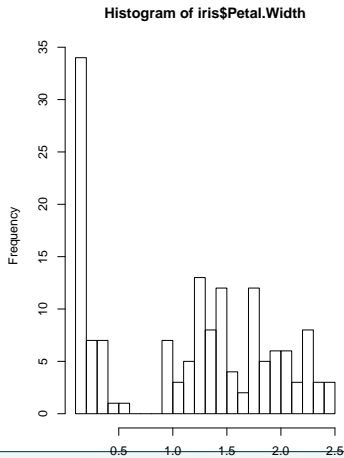
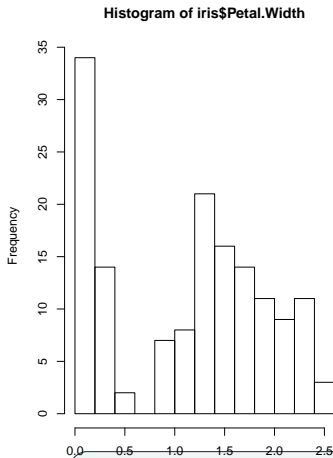


Histograms



► Set the number of breaks to 9, therefore, there will be 10 equal sized bins.

```
par(mfrow=c(1,2))  
hist(iris$Petal.Width, breaks=9)  
hist(iris$Petal.Width, breaks=19)
```




some notes



Note, the number of breaks can be a single number, vector, function to evaluate points or character string.

```
b <- seq(min(iris$Petal.Width), max(iris$Petal.Width),  
         length=11)  
hist(iris$Petal.Width, breaks=b, col="blue", xlab="Petal Width",  
     main="Histogram of Petal Width")
```

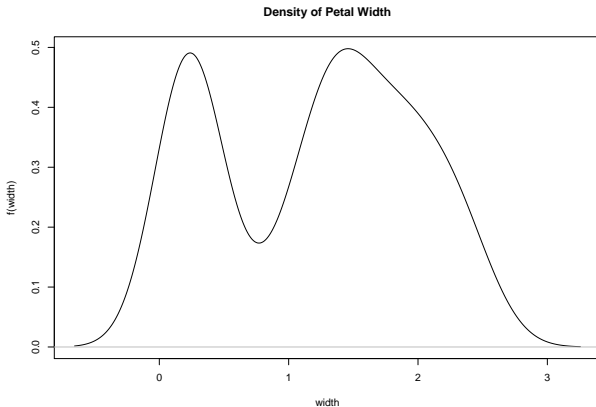




Density Plots

Density plots can be viewed as smoothed versions of a histogram. We can estimate the density using R's density function

```
dens.pw = density(iris$Petal.Width)
plot(dens.pw, ylab = "f(width)", xlab = "width", main= "Density of Petal Width")
```





Density plot cont'd




Let's also look at the density function of `petal.length` for each of the three classes of irises.

```
density.setosa = density(iris$Petal.Length[iris$Species == "setosa"])  
density.versicolor = density(iris$Petal.Length[iris$Species == "versicolor"])  
density.virginica = density(iris$Petal.Length[iris$Species == "virginica"])
```


Density plot cont'd



```
plot(density.setosa, ylab="f(length)", xlab="length (cm)", main="Density plot of  
  Legnths", xlim = c(0,8), lwd=4, col="red")  
lines(density.versicolor, col="blue", lwd=4)  
lines(density.virginica, col="green", lwd=4)  
legend(4, 2.5, c("Setosa", "Versicolor", "Virginica"), lwd=rep(4,3),  
  col=c("red", "blue", "green"))
```



Empirical Cumulative Distribution Functions



The cumulative area under the density function is the cumulative distribution function. Looking at this function using a sample of points gives use the empirical cumulative distribution function (ECDF). Let's start by calculating the cumulative areas


```
ecdf.setosa = ecdf(iris$Petal.Length[iris$Species == "setosa"])
ecdf.versicolor = ecdf(iris$Petal.Length[iris$Species == "versicolor"])
ecdf.virginica = ecdf(iris$Petal.Length[iris$Species == "virginica"])
```

The functions may also be plotted (notice the step function patterns)

```
plot(ecdf.setosa, ylab = "F(length)", xlab = "length (cm)", col.h = "red",
     main = "ECDF for Petal Lengths",
     verticals = T, col.v = "red", do.p = F, lwd = 4, xlim = c(0, 8))
lines(ecdf.versicolor, col.h = "blue", col.v = "blue", verticals = T, lwd = 4,
     do.p = F)
lines(ecdf.virginica, col.h = "green", col.v = "green", verticals = T, lwd = 4,
     do.p = F)
```



ECDF for Petal Lengths



Quantile Plots



The inverse of the ECDF is the quantile function.

```
ps = seq(0,1,length=25)
quantile.setosa = quantile(iris$Petal.Length[iris$Species == "setosa"], probs=ps)
quantile.setosa
```


```
##          0% 4.166667% 8.333333%          12.5% 16.66667% 20.83333%          25%
## 1.000000 1.200000 1.300000 1.300000 1.300000 1.320833 1.400000
## 29.16667% 33.33333%          37.5% 41.66667% 45.83333%          50% 54.16667%
## 1.400000 1.400000 1.400000 1.400000 1.400000 1.500000 1.500000
## 58.33333%          62.5% 66.66667% 70.83333%          75% 79.16667% 83.33333%
## 1.500000 1.500000 1.500000 1.500000 1.575000 1.600000 1.600000
##          87.5% 91.66667% 95.83333%          100%
## 1.600000 1.700000 1.700000 1.900000
```



```
quantile.versicolor = quantile(iris$Petal.Length  
                             [iris$Species == "versicolor"], probs=ps)  
quantile.virginica = quantile(iris$Petal.Length  
                             [iris$Species == "virginica"], probs=ps)  
plot(ps, quantile.setosa, xlab = "p", ylab = "length (cm)",  
     main = "Quantiles for Petal Length",  
     type="l", lwd = 4, col="red", ylim=c(0,8))  
lines(ps, quantile.versicolor, col = "blue", lwd = 4)  
lines(ps, quantile.virginica, col = "green", lwd = 4)
```




Quantiles for Petal Length



check Normality plot



```
hist(iris$Sepal.Length, probability=TRUE, breaks=15, col="light blue")
rug(jitter(iris$Sepal.Length, 5))
points(density(iris$Sepal.Length), type='l', lwd=3, col='red')
f <- function(t) {
  dnorm(t, mean=mean(iris$Sepal.Length), sd=sd(iris$Sepal.Length) )
}
curve(f, add=T, col="blue", lwd=3, lty=2)
```






Q-Q plot

A q-q plot is a plot of the quantiles of the first data set against the quantiles of the second data set.



```
x <- iris$Sepal.Length
qqnorm(x)
qqline(x, col="red", lwd=3)
```





Some more examples

```
par(mfrow=c(1,2))
x <- rnorm(100)
y <- rnorm(100)^2
qqnorm(x, main="Gaussian variable");qqline(x, col="blue", lwd=3)
qqnorm(y, main="Non gaussian variable");qqline(y, col="red", lwd=3)
```






Box Plots

Box plots are used to compactly show many pieces of information about a variables distribution including some summary statistics.


A box plot of each attribute of the iris data set

```
boxplot(iris[,-1])
```



A box plot may also be used for one attribute `petal.length` and show how it varies with the class **Species**

```
boxplot(iris$Petal.Length ~ iris$Species)
```




Scatter Plots



Scatter plots are used to plot two variables against each other (or 3 in the case of 3D plots).

```
plot(iris$Petal.Length, iris$Sepal.Width)
```

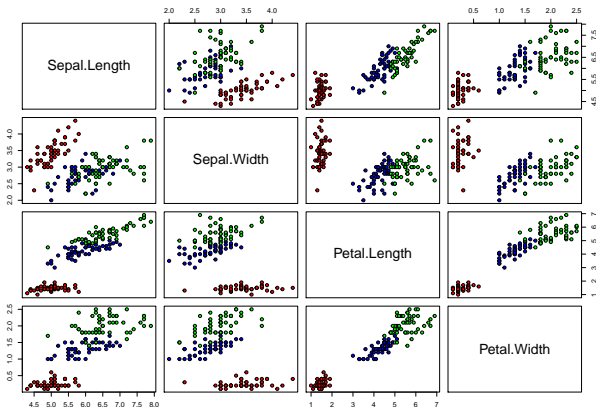


Pairwise Plots



For data sets with only a few attributes, all the pairwise scatter plots may be constructed.

```
pairs(as.matrix(iris[,-5]), pch=21, bg=c("red", "blue", "green")[(iris$Species)])
```



Color Selection Utilities



Default color palette and how to change it

```
palette()
```

```
## [1] "black"    "red"      "green3"   "blue"     "cyan"     "magenta"  "yellow"
## [8] "gray"
```


```
palette(rainbow(5, start=0.1, end=0.2))
palette()
```

```
## [1] "#FF9900" "#FFBF00" "#FFE600" "#F2FF00" "#CCFF00"
```



The gray function allows to select any type of gray shades by providing values from 0 to 1



```
pie(1:5,col=gray(seq(0.1, 1, by= 0.2)))
```



Color gradients with colorpanel function from gplots library



```
library(gplots, warn.conflicts = F)
par(mfrow=c(1,2))
pie(1:50, col=colorpanel(50,"red","black","green"))
pie(1:50, col=rainbow(50))
```






Arranging Plots with Variable Width




The layout function allows to divide the plotting device into variable numbers of rows and columns with the column-widths and the row-heights specified in the respective arguments.

```
nf <- layout(matrix(c(2,1,3,3), 2, 2, byrow=TRUE), c(3,7), c(5,3))
layout.show(nf)
```





```
x <- rnorm(100) ; y <- rnorm(100)
layout(matrix(c(2,1,3,3), 2, 2, byrow=TRUE), c(3,7), c(5,4))
plot(x,y)
boxplot(y,horizontal = FALSE,col=2)
hist(x,col=4)
```






Margine plots



```
data(cars)
layout( matrix( c(2,1,0,3), 2, 2, byrow=T ),
c(1,6), c(4,1))
par(mar=c(1,1,5,2)); plot(cars$dist ~ cars$speed,xlab='', ylab='',las = 1)
rug(side=1, jitter(cars$speed, 5) );rug(side=2, jitter(cars$dist, 20) );
title(main = "cars data")
par(mar=c(1,2,5,1))
boxplot(cars$dist, axes=F); title(ylab='Stopping distance (ft)', line=0)
par(mar=c(5,1,1,2))
boxplot(cars$speed, horizontal=T, axes=F); title(xlab='Speed (mph)', line=1)
```





Exercise: Bar Plots, Pie plot,



Consider the same approach in this slides for `barplot`



Saving Graphics to Files



After the `pdf()` command all graphs are redirected to file `test.pdf`. Works for all common formats similarly: `jpeg`, `png`, `ps`, `tiff`, ...


```
pdf("test.pdf"); plot(1:10, 1:10); dev.off()
```

3 Dimesional functions and their plots

persp 3d plot



```
x <- seq(-10, 10, length= 30)
y <- x
f <- function(x, y) { r <- sqrt(x^2+y^2); 10 * sin(r)/r }
z <- outer(x, y, f)
persp(x, y, z, theta = 30, phi = 30, expand = 0.5, col = "lightblue")
```






Countour Plots



A contour plot is a graphical technique for representing a 3-dimensional surface by plotting constant z slices, called contours, on a 2-dimensional format.

```
contour(x,y,z) # or contour(z)
```




2D Density



Density estimation is available for higher dimensional data. Use the **MASS** package.


```
library(MASS)
petal.dens = kde2d(iris$Petal.Length, iris$Petal.Width)
contour(petal.dens)
```





The plot may aslo be viewed as an image

```
image(petal.dens)
```



The background of the slide is a light gray grid of triangles and squares. The triangles are arranged in a pattern that creates a sense of depth and movement, with some triangles pointing upwards and others downwards. The squares are interspersed among the triangles, forming a larger grid structure.

lattice Grapphics

What is lattice?




- ▶ High-level graphics system
- ▶ Implements Trellis graphics system from S-Plus Simplifies high-level plotting tasks: arranging complex graphical features
- ▶ Syntax similar to R's base graphics

Scatter Plot Sample




```
library(lattice)  
xyplot(1:8 ~ 1:8 | rep(LETTERS[1:4], each=2), as.table=TRUE)
```



Line Plot Sample



```
parallelplot(~iris[1:4] | Species, iris, horizontal.axis = FALSE,  
            layout = c(1, 3, 1))
```



ggplot2 package (a superfull package in graphic)

- ▶ What is ggplot2?
 - ▶ High-level graphics system
 - ▶ Streamlines many graphics workflows for complex plots
 - ▶ Simpler qplot function provides many shortcuts
- ▶ ggplot function accepts two arguments
 - ▶ Data set to be plotted
 - ▶ **Aesthetic** mappings provided by **aes** function
 - ▶ Additional parameters such as geometric objects (e.g. points, lines, bars) are passed on by appending them with + as separator.

qplot function



qplot syntax is similar to R's basic plot function Arguments:

- ▶ x: x-coordinates (e.g. col1)
- ▶ y: y-coordinates (e.g. col2)
- ▶ data: data frame with corresponding column names xlim, ylim:
e.g. xlim=c(0,10)
- ▶ log: e.g. log="x" or log="xy"
- ▶ main: main title; see ?plotmath for mathematical formula xlab, ylab: labels for
the x- and y-axes
- ▶ color, shape, size
- ▶ ...: many arguments accepted by plot function




Create sample data

```
library(ggplot2,quietly = T)
```

```
## Warning: package 'ggplot2' was built under R version 3.2.5
```


```
#cat <- rep(c("A", "B"), 5)  
qplot(iris$Sepal.Length, iris$Sepal.Width, geom="point")
```



Prints dots with different sizes and colors




```
ggplot(Sepal.Length, Petal.Width, data = iris, geom="point",
       size=Petal.Length,col=Species,
       main="Dot Size and Color Relative to Some Values")
```



Drops legend




```
ggplot(Sepal.Length, Petal.Width, data = iris, geom="point", col=Species)+
  theme(legend.position = "none")
```



Plot different shapes




```
ggplot(Sepal.Length, Petal.Width, data = iris, geom="point",
       shape=iris$Species, size=2)
```



Scatter Plot with Regression Line



```
ggplot(Sepal.Length, Petal.Width, data = iris, geom = c("point"))+  
  geom_smooth(method=lm,col=2)
```



More application of qplot



```
qplot(Sepal.Length, Petal.Width, data = iris, facets = Species ~ Petal.Length)
```

