



Generating Random Number and Distributions as Models

S. Morteza Najibi, Shiraz University

April 24, 2019

- ▶ How *does* R get “random” numbers, anyway?
- ▶ It doesn't, really – it uses a trick that should be indistinguishable from the real McCoy

Pseudorandom generators produce a deterministic sequence that is indistinguishable from a true random sequence if you don't know how it started.

```
runif(1:10)
```

```
## [1] 0.263972099 0.227739427 0.632124108 0.958168537 0.009046437  
## [6] 0.622218946 0.882597443 0.765022868 0.451999373 0.902534678
```

```
set.seed(10)  
runif(1:10)
```

```
## [1] 0.50747820 0.30676851 0.42690767 0.69310208 0.08513597 0.22543662  
## [7] 0.27453052 0.27230507 0.61582931 0.42967153
```

```
set.seed(10)  
runif(1:10)
```

```
## [1] 0.50747820 0.30676851 0.42690767 0.69310208 0.08513597 0.22543662  
## [7] 0.27453052 0.27230507 0.61582931 0.42967153
```

A few distributions of interest:

- ▶ Discret Uniform $\{1, \dots, n\}$
- ▶ Uniform(0,1)
- ▶ Bernoulli(p)
- ▶ Binomial(n,p)
- ▶ normal(mu,sigma)
- ▶ Exponential(lambda)
- ▶ Gamma(n,lambda)

Use `sample` function to generate random number from a discrete uniform distribution

```
sample(1:10,size = 5,replace = T)
```

```
## [1] 7 6 2 6 4
```

```
sample(10,size = 5,replace = T) # if the first argument is not vector
```

```
## [1] 5 1 3 4 9
```

Suppose we were working with the Exponential distribution.

- ▶ `rexp()` generates variates from the distribution.
- ▶ `dexp()` gives the probability density function.
- ▶ `pexp()` gives the cumulative distribution function.
- ▶ `qexp()` gives the quantiles.

- ▷ `dfoo` = the probability *d*ensity (if continuous) or probability mass function of *foo* (pdf or pmf)
- ▷ `pfoo` = the cumulative *p*robability function (CDF)
- ▷ `qfoo` = the *q*uantile function (inverse to CDF)
- ▷ `rfoo` = draw *r*andom numbers from *foo* (first argument always the number of draws)

?Distributions to see which distributions are built in

If you write your own, follow the conventions



```
dnorm(x=c(-1,0,1),mean=1,sd=0.1)
```

```
## [1] 5.520948e-87 7.694599e-22 3.989423e+00
```

```
pnorm(q=c(2,-2)) # defaults to mean=0,sd=1
```

```
## [1] 0.97724987 0.02275013
```

```
dbinom(5,size=7,p=0.7,log=TRUE)
```

```
## [1] -1.146798
```

```
qchisq(p=0.95,df=5)
```

```
## [1] 11.0705
```

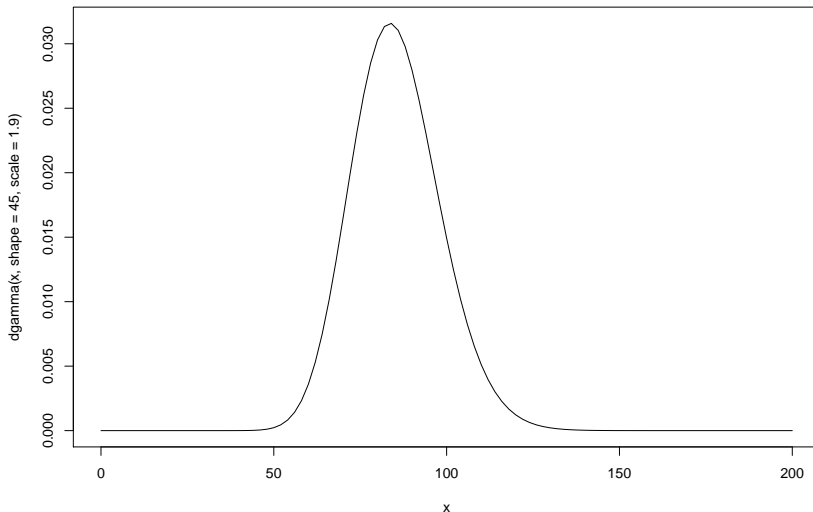
```
rt(n=4,df=2)
```

```
## [1] 1.0154673 -0.1983909 0.9861791 -1.5666848
```

Displaying Probability Distributions

curve is very useful for the d, p, q functions:

```
curve(dgamma(x,shape=45,scale=1.9),from=0,to=200)
```



- ▶ Match moments (mean, variance, etc.)
- ▶ Match other summary statistics
- ▶ Maximize the likelihood

- ▶ Pick enough moments that they **identify** the parameters
 - ▶ At least 1 moment per parameter; algebraically independent

- ▶ Write equations for the moments in terms of the parameters
e.g., for gamma

$$E(X) = \bar{x} , E(X^2) = \bar{x}^2$$

- ▶ Do the algebra by hand to solve the equations

$$shape = \bar{x}^2 / s^2 , scale = s^2 / \bar{x}$$

```
gamma.est_MM <- function(x) {  
  m <- mean(x); v <- var(x)  
  return(c(shape=m^2/v, scale=v/m))  
}
```

- ▶ Usually we think of the parameters as fixed and consider the probability of different outcomes, $f(x; \theta)$ with θ constant and x changing **Likelihood** of a parameter
- ▶ With independent data points x_1, x_2, \dots, x_n , likelihood is

$$L(\theta) = \prod_{i=1}^n f(x_i; \theta)$$

- ▶ Multiplying lots of small numbers is numerically bad; take the log:

$$\ell(\theta) = \sum_{i=1}^n \log f(x_i; \theta)$$

▶ In pseudo-code:

```
loglike.foo <- function(params, x) {  
  sum(dfoo(x=x,params,log=TRUE))  
}
```

- ▶ We maximize it!
- ▶ Sometimes we can do the maximization by hand with some calculus
 - ▶ For Gaussian, MLE = just match the mean and variance
 - ▶ For Pareto, $\text{MLE } \hat{a} = 1 + 1/\log(x/x_{\min})$
- ▶ Doing numerical optimization
 - ▶ Stick in a minus sign if we're using a minimization function

MLE for one-dimensional distributions can be done through `fitdistr` in the **MASS** package

It knows about most the standard distributions, but you can also give it arbitrary probability density functions and it will try to maximize them

A starting value for the optimization is optional for some distributions, required for others (including user-defined densities)

Returns the parameter estimates and standard errors
SEs come from large- n approximations so use cautiously

Fit the gamma distribution to the cats' hearts:

```
require(MASS)
```

```
## Loading required package: MASS
```

```
fitdistr(cats$Hwt, densfun="gamma")
```

```
##      shape      rate  
## 20.2998092 1.9095724  
## ( 2.3729250) ( 0.2259942)
```

Returns: estimates above, standard errors below

▶ simulate, then estimate; estimates should converge as the sample grows

```
gamma.est_MM(rgamma(100,shape=19,scale=45))
```

```
##      shape      scale  
## 20.03284 43.16644
```

```
gamma.est_MM(rgamma(1e5,shape=19,scale=45))
```

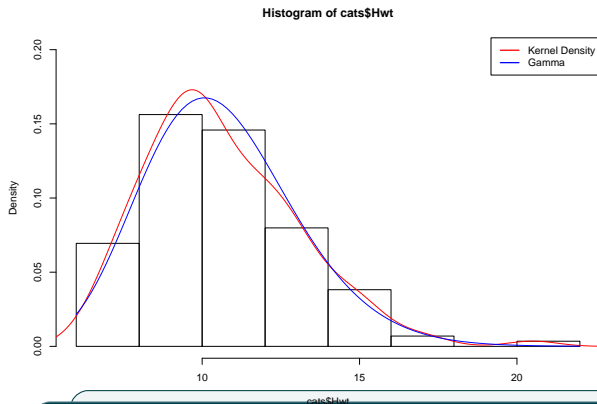
```
##      shape      scale  
## 19.00715 44.88471
```

```
gamma.est_MM(rgamma(1e6,shape=19,scale=45))
```

```
##      shape      scale  
## 18.99766 45.01638
```

Use your eyes: Graphic overlays of theory vs. data

```
hist(cats$Hwt,prob=T,ylim=c(0,.2))  
lines(density(cats$Hwt),col=2)  
cats.gamma <- gamma.est_MM(cats$Hwt)  
curve(dgamma(x,shape=cats.gamma["shape"],scale=cats.gamma["scale"]),add=TRUE,col="red",lty=1,lwd=2)  
legend("topright",c("Kernel Density","Gamma"),col=c(2,4),lty=1,lwd=2)
```



- ▶ How much should the QQ plot wiggle around the diagonal?
- ▶ Answer a different question...
- ▶ Biggest gap between theoretical and empirical CDF:

$$D_{KS} = \max_x |F(x) - \hat{F}(x)|$$

- ▶ Useful because D_{KS} always has the same distribution *if* the theoretical CDF is fixed and correct
- ▶ Also works for comparing the empirical CDFs of two samples, to see if they came from the same distribution

```
test.data <- rnorm(100,.5,.1)
ks.test(test.data,pnorm,mean=.1,sd=0.1)
```

```
##
## One-sample Kolmogorov-Smirnov test
##
## data: test.data
## D = 0.97792, p-value < 2.2e-16
## alternative hypothesis: two-sided
```

Ex: How does it works for other distributions?

Compare an actual table of counts to a hypothesized probability distribution:

```
coin <- rbinom(100,1,.45); chisq.test(table(coin),p=c(1/2,1/2))
```

```
##  
## Chi-squared test for given probabilities  
##  
## data:  table(coin)  
## X-squared = 0.04, df = 1, p-value = 0.8415
```

```
coin <- rbinom(1000,1,.45) ; chisq.test(table(coin),p=c(1/2,1/2))
```

```
##  
## Chi-squared test for given probabilities  
##  
## data:  table(coin)  
## X-squared = 29.584, df = 1, p-value = 5.355e-08
```

- ▶ The df is the number of cells in the table -1
- ▶ If we estimate q parameters, we need to subtract q degrees of freedom

- ▶ Divide the range into bins and count the number of observations in each bin; this will be `x` in `chisq.test()`
- ▶ Use the CDF function `p.foo` to calculate the theoretical probability of each bin; this is `p`
- ▶ Plug in to `chisq.test`
- ▶ If parameters are estimated, adjust

▷ `hist()` gives us break points and counts:

```
cats.hist <- hist(cats$Hwt,plot=FALSE)
cats.hist$breaks
```

```
## [1] 6 8 10 12 14 16 18 20 22
```

```
cats.hist$counts
```

```
## [1] 20 45 42 23 11 2 0 1
```

Use these for a χ^2 test:

```
# Why the padding by -Inf and Inf?  
p <- diff(pgamma(c(-Inf,cats.hist$breaks,Inf),shape=cats.gamma["shape"],  
                 scale=cats.gamma["scale"])))  
# Why the padding by 0 and 0?  
chisq.test(c(0,cats.hist$counts,0),p=p)
```

```
##  
## Chi-squared test for given probabilities  
##  
## data:  c(0, cats.hist$counts, 0)  
## X-squared = 12.133, df = 9, p-value = 0.2059
```

Don't need to run `hist` first; can also use `cut` to discretize (see `?cut`)

- ▶ Visualizing and computing empirical distribution
- ▶ Parametric distributions are models
- ▶ Methods of fitting: moments and likelihood
- ▶ Methods of checking: visual comparisons, other statistics, tests, calibration